

face-recognition-by-transfer-learning

In this project, I have created a Face-Recognition model using the concept of **Feature Tuning**.

Step 1: We start off by collecting our dataset. For this, I have used *Haarcascade FrontalFace*. I have collected 200 images of mine & my friend for training the model & 100 images each for testing the model. You can use the following code to collect the images and prepare the dataset.

```
#MY TRAIN IMAGES
import cv2
import numpy as np

# Load HAAR face classifier
face_classifier = cv2.CascadeClassifier('haarcascade_frontalface_default.xml')

# Load functions
def face_extractor(img):
    # Function detects faces and returns the cropped face
    # If no face detected, it returns the input image

    #gray = cv2.cvtColor(img,cv2.COLOR_BGR2GRAY)
    faces = face_classifier.detectMultiScale(img, 1.3, 5)

    if faces is ():
        return None

    # Crop all faces found
    for (x,y,w,h) in faces:
        cropped_face = img[y:y+h, x:x+w]

    return cropped_face

# Initialize Webcam
cap = cv2.VideoCapture(0)
count = 0

# Collect 100 samples of your face from webcam input
while True:

    ret, frame = cap.read()
    if face_extractor(frame) is not None:
        count += 1
        face = cv2.resize(face_extractor(frame), (224, 224))
        #face = cv2.cvtColor(face, cv2.COLOR_BGR2GRAY)

        # Save file in specified directory with unique name
        file_name_path = 'C://Users//Anuddeeph Nalla//Desktop//MLOPS-WS//dataset//train//Anudeep//anudeep' + str(count) + '.jpg'
        cv2.imwrite(file_name_path, face)

        # Put count on images and display live count
        cv2.putText(face, str(count), (50, 50), cv2.FONT_HERSHEY_COMPLEX, 1, (0,255,0), 2)
        cv2.imshow('Face Cropper', face)

    else:
        print("Face not found")
        pass

    if cv2.waitKey(1) == 13 or count == 200: #13 is the Enter Key
        break

cap.release()
cv2.destroyAllWindows()
print("Collecting Samples Complete")
```

I have used the same block of code multiple times to collect all the training & testing images of me & my friend. You can collect images of more people as per your requirement. For more details check the code

Step 2: Now, we import pre-created MobileNet model from keras.applications. We freeze the already trained layers by *layer.trainable= False*.

```
In [5]: from keras.applications import MobileNet

# MobileNet was designed to work on 224 x 224 pixel input images sizes
img_rows, img_cols = 224, 224

# Re-Loads the MobileNet model without the top or FC layers
MobileNet = MobileNet(weights = 'imagenet',
                        include_top = False,
                        input_shape = (img_rows, img_cols, 3))

# Here we freeze the last 4 layers
# Layers are set to trainable as True by default
for layer in MobileNet.layers:
    layer.trainable = False

# Let's print our layers
for (i, layer) in enumerate(MobileNet.layers):
    print(str(i) + " " + layer.__class__.__name__ + ", layer.trainable")

Using TensorFlow backend.

0 InputLayer False
1 ZeroPadding2D False
2 Conv2D False
3 BatchNormalization False
4 ReLU False
5 DepthwiseConv2D False
6 BatchNormalization False
7 ReLU False
8 Conv2D False
9 BatchNormalization False
10 ReLU False
11 ZeroPadding2D False
12 DepthwiseConv2D False
13 BatchNormalization False
14 ReLU False
15 Conv2D False
16 BatchNormalization False
17 ReLU False
18 DepthwiseConv2D False
19 BatchNormalization False
20 ReLU False
21 Conv2D False
22 BatchNormalization False
23 ReLU False
24 ZeroPadding2D False
25 DepthwiseConv2D False
26 BatchNormalization False
27 ReLU False
28 Conv2D False
29 BatchNormalization False
30 ReLU False
31 DepthwiseConv2D False
32 BatchNormalization False
33 ReLU False
34 Conv2D False
35 BatchNormalization False
36 ReLU False
```

```
In [8]: from keras_preprocessing.image import ImageDataGenerator
```

```
In [10]: train_datagen = ImageDataGenerator(
        rescale=1./255,
        shear_range=0.2,
        zoom_range=0.2,
        horizontal_flip=True)
test_datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=45,
    width_shift_range=0.3,
    height_shift_range=0.3,
    horizontal_flip=True,
    fill_mode='nearest')
train_generator = train_datagen.flow_from_directory(
    'dataset/train',
    target_size=(img_rows, img_cols),
    batch_size=64,
    class_mode='categorical')
validation_generator = test_datagen.flow_from_directory(
    'dataset/test',
    target_size=(img_rows, img_cols),
    batch_size=32,
    class_mode='categorical')
```

```
Found 400 images belonging to 2 classes.
Found 200 images belonging to 2 classes.
```

Step 3: We add layers as per our requirement. Here, I have used *Softmax* activation function.

```
In [6]: def lw(bottom_model, num_classes):
        """creates the top or head of the model that will be
        placed ontop of the bottom layers"""

        top_model = bottom_model.output
        top_model = GlobalAveragePooling2D()(top_model)
        top_model = Dense(1024,activation='relu')(top_model)
        top_model = Dense(1024,activation='relu')(top_model)
        top_model = Dense(512,activation='relu')(top_model)
        top_model = Dense(num_classes,activation='softmax')(top_model)
        return top_model
```

```
In [7]: from keras.models import Sequential
        from keras.layers import Dense, Dropout, Activation, Flatten, GlobalAveragePooling2D
        from keras.layers import Conv2D, MaxPooling2D, ZeroPadding2D
        from keras.layers.normalization import BatchNormalization
        from keras.models import Model

        # Set our class number to 3 (Young, Middle, Old)
        num_classes = 2

        FC_Head = lw(MobileNet, num_classes)

        model = Model(inputs = MobileNet.input, outputs = FC_Head)

        print(model.summary())
```

conv_dw_12 (DepthwiseConv2D)	(None, 7, 7, 512)	4608
conv_dw_12_bn (BatchNormaliz	(None, 7, 7, 512)	2048
conv_dw_12_relu (ReLU)	(None, 7, 7, 512)	0
conv_pw_12 (Conv2D)	(None, 7, 7, 1024)	524288
conv_pw_12_bn (BatchNormaliz	(None, 7, 7, 1024)	4096
conv_pw_12_relu (ReLU)	(None, 7, 7, 1024)	0
conv_dw_13 (DepthwiseConv2D)	(None, 7, 7, 1024)	9216
conv_dw_13_bn (BatchNormaliz	(None, 7, 7, 1024)	4096
conv_dw_13_relu (ReLU)	(None, 7, 7, 1024)	0
conv_pw_13 (Conv2D)	(None, 7, 7, 1024)	1048576

Next, we load our dataset. We have used the augmentation technique to increase our dataset since the size of original dataset is too small for a good accuracy.

Step 5: Now, we begin training our model.

```
In [11]: from keras.optimizers import RMSprop
from keras.callbacks import ModelCheckpoint, EarlyStopping
from keras.models import load_model

checkpoint = ModelCheckpoint("facer.h5",
                             monitor="val_loss",
                             mode="min",
                             save_best_only = True,
                             verbose=1)

earlystop = EarlyStopping(monitor = 'val_loss',
                           min_delta = 0,
                           patience = 3,
                           verbose = 1,
                           restore_best_weights = True)

# we put our call backs into a callback list
callbacks = [earlystop, checkpoint]

# We use a very small learning rate
model.compile(loss = 'categorical_crossentropy',
              optimizer = RMSprop(lr = 0.001),
              metrics = ['accuracy'])
```

```
In [12]: # Enter the number of training and validation samples here
nb_train_samples = 544
nb_validation_samples = 200

# We only train 10 EPOCHS
epochs = 5
batch_size = 64

history = model.fit_generator(
    train_generator,
    steps_per_epoch = nb_train_samples // batch_size,
    epochs = epochs,
    callbacks = callbacks,
    validation_data = validation_generator,
    validation_steps = nb_validation_samples // batch_size)

classifier = load_model('facer.h5')

Epoch 1/5
8/8 [=====] - 115s 14s/step - loss: 4.1117 - accuracy: 0.5797 - val_loss: 0.0930 - val_accuracy: 0.9896

Epoch 00001: val_loss improved from inf to 0.09295, saving model to facer.h5
Epoch 2/5
8/8 [=====] - 98s 12s/step - loss: 0.0152 - accuracy: 0.9914 - val_loss: 0.0980 - val_accuracy: 0.9896
```

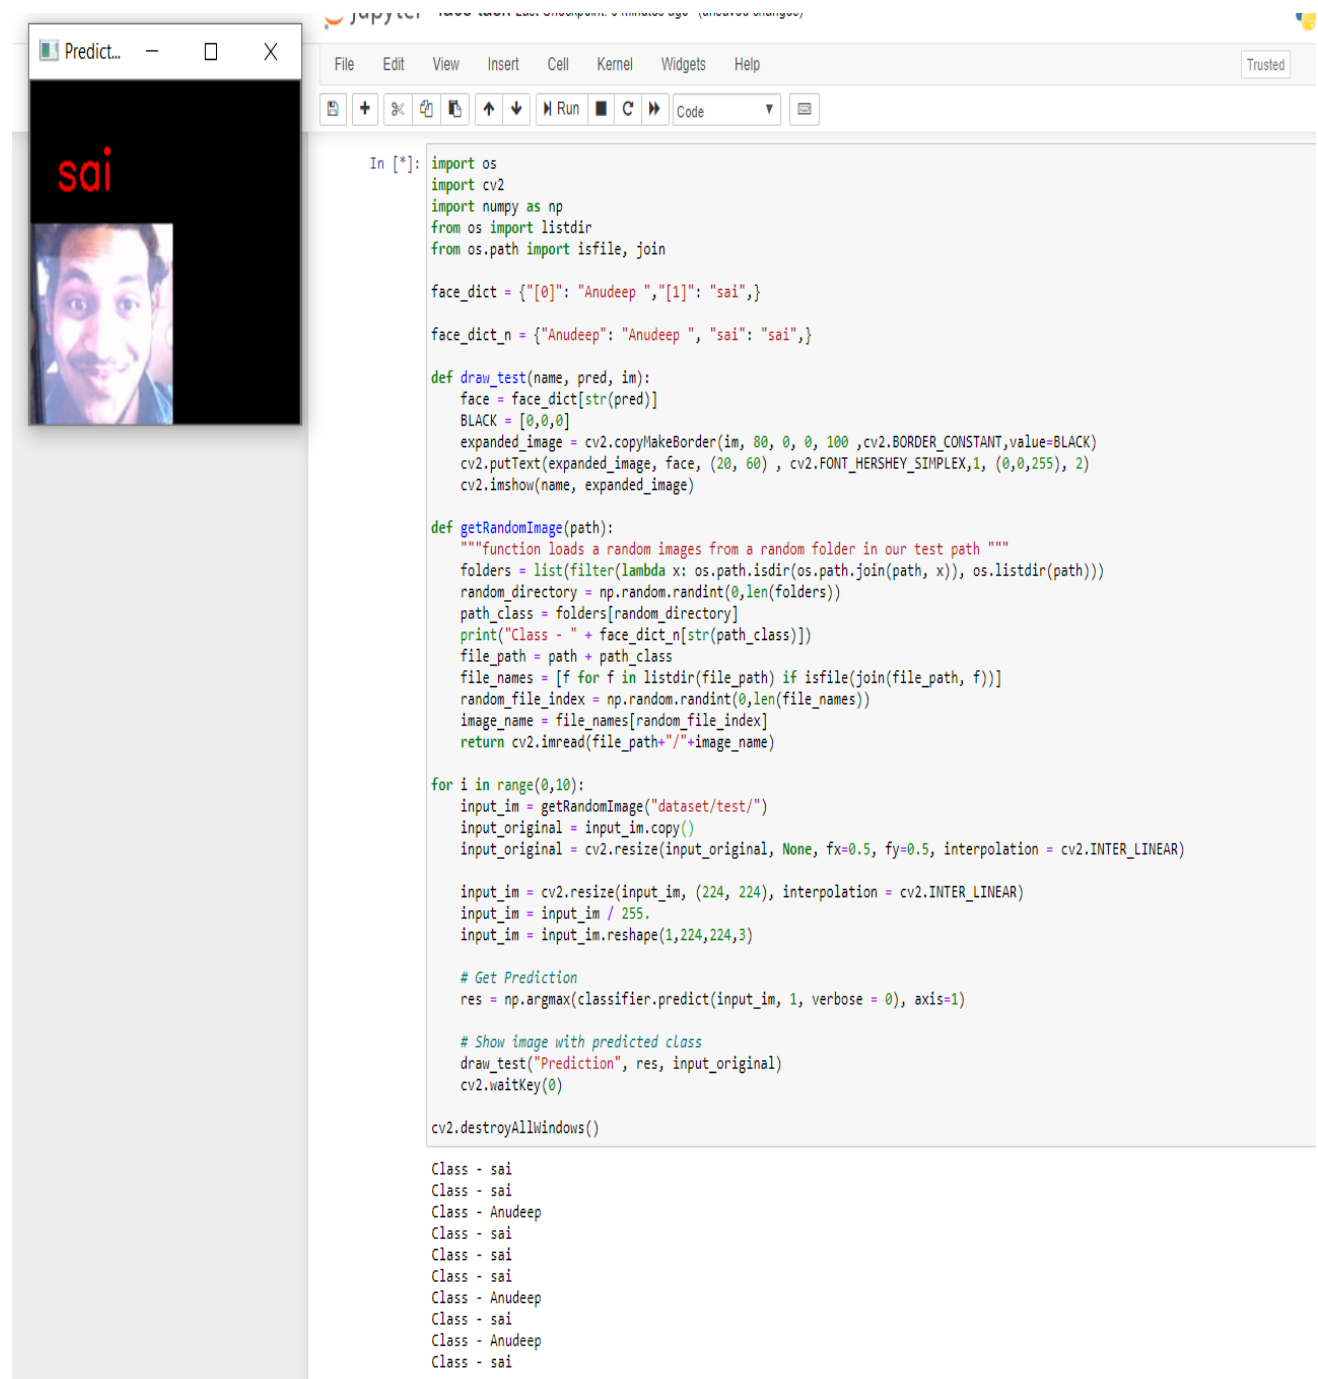
The model has been effectively trained and ready to use. You can use this model for prediction.

In this model, I got 99% accuracy, because the data was very less.

Step 6: Now, I have loaded the created model for prediction, and predicted mine and my friend face.

```
In [13]: from keras.models import load_model  
         classifier = load_model('facer.h5')
```

The output of predicted model:



```
In [*]: import os  
import cv2  
import numpy as np  
from os import listdir  
from os.path import isfile, join  
  
face_dict = {"[0]": "Anudeep ", "[1]": "sai"},  
  
face_dict_n = {"Anudeep": "Anudeep ", "sai": "sai"},  
  
def draw_test(name, pred, im):  
    face = face_dict[str(pred)]  
    BLACK = [0,0,0]  
    expanded_image = cv2.copyMakeBorder(im, 80, 0, 0, 100 ,cv2.BORDER_CONSTANT,value=BLACK)  
    cv2.putText(expanded_image, face, (20, 60) , cv2.FONT_HERSHEY_SIMPLEX,1, (0,0,255), 2)  
    cv2.imshow(name, expanded_image)  
  
def getRandomImage(path):  
    """function loads a random images from a random folder in our test path """  
    folders = list(filter(lambda x: os.path.isdir(os.path.join(path, x)), os.listdir(path)))  
    random_directory = np.random.randint(0,len(folders))  
    path_class = folders[random_directory]  
    print("Class - " + face_dict_n[str(path_class)])  
    file_path = path + path_class  
    file_names = [f for f in listdir(file_path) if isfile(join(file_path, f))]  
    random_file_index = np.random.randint(0,len(file_names))  
    image_name = file_names[random_file_index]  
    return cv2.imread(file_path+"/"+image_name)  
  
for i in range(0,10):  
    input_im = getRandomImage("dataset/test/")  
    input_original = input_im.copy()  
    input_original = cv2.resize(input_original, None, fx=0.5, fy=0.5, interpolation = cv2.INTER_LINEAR)  
  
    input_im = cv2.resize(input_im, (224, 224), interpolation = cv2.INTER_LINEAR)  
    input_im = input_im / 255.  
    input_im = input_im.reshape(1,224,224,3)  
  
    # Get Prediction  
    res = np.argmax(classifier.predict(input_im, 1, verbose = 0), axis=1)  
  
    # Show image with predicted class  
    draw_test("Prediction", res, input_original)  
    cv2.waitKey(0)  
  
cv2.destroyAllWindows()  
  
Class - sai  
Class - sai  
Class - Anudeep  
Class - sai  
Class - sai  
Class - sai  
Class - Anudeep  
Class - sai  
Class - Anudeep  
Class - sai
```